Lab Exercise 3- Creating and Managing a ReplicaSet in Kubernetes

Objective:

A ReplicaSet in Kubernetes ensures a specified number of Pod replicas are running at any given time. This exercise will guide you through creating a ReplicaSet to maintain the desired state of your application.

- Understand the syntax and structure of a Kubernetes ReplicaSet definition file (YAML).
- Learn how to create and manage a ReplicaSet to ensure application availability.
- Understand how a ReplicaSet helps in scaling applications and maintaining desired states.

Prerequisites

- Kubernetes Cluster: Have a running Kubernetes cluster (locally using Minikube or kind, or a cloud-based service).
- kubectl: Install and configure kubectl to interact with your Kubernetes cluster.
- Basic Knowledge of YAML: Familiarity with YAML format will be helpful for understanding Kubernetes resource definitions.

Step-by-Step Guide

Step 1: Understanding ReplicaSet

A ReplicaSet ensures a specified number of Pod replicas are running at any given time. If a Pod crashes or is deleted, the ReplicaSet creates a new one to meet the defined number of replicas. This helps maintain application availability and ensures that your application can handle increased load by distributing traffic among multiple Pods.

Step 2: Create a ReplicaSet

We'll define a ReplicaSet to maintain three replicas of a simple Nginx web server Pod. Create a YAML file named **replicaset.yaml** with the following content:

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: nginx-replicaset
spec:
 replicas: 3
 selector:
  matchLabels:
   app: my-app
 template:
  metadata:
   labels:
    app: my-app
  spec:
   containers:
   - name: nginx
     image: nginx:latest
     ports:
     - containerPort: 80
```

Explanation:

- apiVersion: Defines the API version (apps/v1) used for the ReplicaSet resource.
- kind: Specifies that this resource is a ReplicaSet.
- metadata: Contains metadata about the ReplicaSet, including name.

- o   name: The unique name for the ReplicaSet.
- spec: Provides the specification for the ReplicaSet.
    - o   replicas: Defines the desired number of Pod replicas.
    - o   selector: Criteria for selecting Pods managed by this ReplicaSet.
        - ▪   matchLabels: Labels that Pods must have to be managed by this ReplicaSet.
    - o   template: Defines the Pod template used for creating new Pods.
        - ▪   metadata: Contains metadata for the Pods, including labels.
            - labels: Labels applied to Pods created by this ReplicaSet.
    - o   spec: Specification for the Pods.
        - ▪   containers: Lists the containers that will run in the Pod.
            - name: The unique name of the container within the Pod.
            - image: The Docker image used for the container.
            - ports: Ports exposed by the container.

Step 3: Apply the YAML to Create the ReplicaSet

Use the kubectl apply command to create the ReplicaSet based on the YAML file.

```
kubectl apply -f replicaset.yaml
```

Verify the ReplicaSet is running and maintaining the desired number of replicas:

```
kubectl get replicaset
```

This command lists all ReplicaSets in the current namespace.

To check the Pods created by the ReplicaSet:

```
kubectl get pods -l app=my-app
```

This command lists all Pods with the label app=my-app.

Step 4: Managing the ReplicaSet

1. Scaling the ReplicaSet

You can scale the number of replicas managed by the ReplicaSet using the kubectl scale command.

```
kubectl scale --replicas=5 replicaset/nginx-replicaset
```

This command scales the ReplicaSet to maintain 5 replicas. Verify the scaling operation:

```
kubectl get pods -l app=my-app
```

You should see that the number of Pods has increased to 5.

2. Updating the ReplicaSet

If you need to update the Pod template (e.g., to use a different Docker image version), modify the YAML file and apply it again. For instance, change the image to a specific version of Nginx:

```
spec:
 template:
  spec:
   containers:
   - name: nginx
     image: nginx:1.19.3  # Change to a specific version
```

Apply the changes:

```
kubectl apply -f nginx-replicaset.yaml
```

Check the status to ensure the Pods are updated:

```
kubectl get pods -l app=nginx
```

Note: Updating a ReplicaSet doesn't automatically replace existing Pods with new ones. In practice, you often create a new ReplicaSet or Deployment for updates.

3. Deleting the ReplicaSet

To clean up the ReplicaSet and its Pods, use the kubectl delete command:

```
kubectl delete -f nginx-replicaset.yaml
```

This command deletes the ReplicaSet and all the Pods managed by it.